

Apache Cassandra

“manage massive amounts of data, fast, without losing sleep.”



Fernando Fagonde
UNIPAMPA
SABD
Bagé - RS

Summary

- What is Cassandra?
 - NoSQL*;
 - CAP Theorem;
 - Proven;
 - Fault Tolerant;
 - Performant;
 - Decentralized;
 - Scalable;
 - Durable;
 - Elastic;
 - CQL.
- Installation.
- Next Steps.

What is Cassandra?



+



=



What is Cassandra?

- NoSQL*

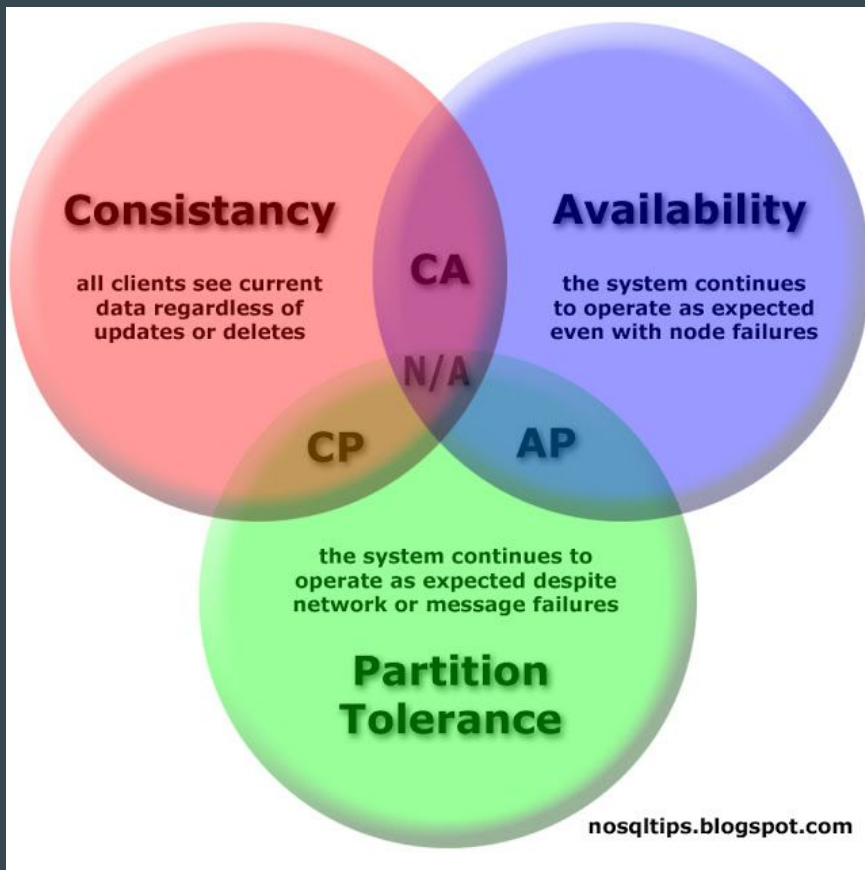
*NoSQL

- NoSQL é um termo usado para descrever bancos de dados não relacionais de alto desempenho.
- Os bancos de dados NoSQL usam diversos modelos de dados, incluindo documentos, gráficos, chave-valor e colunares.
- Bancos de dados NoSQL são amplamente reconhecidos pela facilidade de desenvolvimento, desempenho escalável, alta disponibilidade e resiliência.

Fonte: <https://aws.amazon.com/pt/nosql/>

CAP Theorem

Dr. Eric Brewer



CAP Theorem

“Cassandra are always weakly consistent, favoring high availability and response time (latency), ie PA / L.”

fonte:<http://blog.caelum.com.br/nosql-do-teorema-cap-para-paccl/>

Proven

- Cassandra is in use at enterprises that have large, active data sets:
 - CERN;
 - eBay;
 - Github;
 - GoDaddy;
 - Instagram;
 - Netflix;
 - Reddit;
 - The Weather Channel;
 - +1500.

fontes:

<http://cassandra.apache.org/>

Fault Tolerant

- Data is automatically replicated to multiple nodes for fault-tolerance;
- Replication across multiple data centers is supported;
- Failed nodes can be replaced with no downtime.

fontes:

<http://cassandra.apache.org/>

Performant

- Cassandra Consistently outperforms popular NoSQL alternatives in benchmarks and real applications, primarily because of fundamental architectural choices;

fontes:

<http://cassandra.apache.org/>

Decentralized

- There are no single points of failure.
- There are no network bottlenecks.
- Every node in the cluster is identical.

fontes:

<http://cassandra.apache.org/>

Scalable

- Some of the largest production deployments include:
 - Apple's, with over 75,000 nodes storing over 10 PB of data;
 - Netflix (2,500 nodes, 420 TB, over 1 trillion requests per day);
 - Chinese search engine Easou (270 nodes, 300 TB, over 800 million requests per day);
 - eBay (over 100 nodes, 250 TB).

fontes:

<http://cassandra.apache.org/>

Durable

- Cassandra is suitable for applications that can't afford to lose data, even when an entire data center goes down;

fontes:

<http://cassandra.apache.org/>

Elastic

- Read and write throughput both increase linearly as new machines are added, with no downtime or interruption to applications.

fontes:

<http://cassandra.apache.org/>

CQL

- **Cassandra Query Language (CQL)**

fontes:

<http://cassandra.apache.org/>

Installation

- Debian like:
 - `echo "deb http://www.apache.org/dist/cassandra/debian 310x main" | sudo tee -a /etc/apt/sources.list.d/cassandra.sources.list`
 - `sudo apt-get install cassandra`
- Docker
 - `sudo docker run --name instance-name -d cassandra:3.10 (stable)`

fontes:

https://hub.docker.com/_/cassandra/

<http://cassandra.apache.org/download/>

Next Steps

- Testar com uma aplicação.

A person with long brown hair, wearing a red long-sleeved shirt, is holding a white rectangular sign with both hands. The sign has the word 'Obrigado' written on it in a black, rounded, handwritten-style font. The background is solid black.

Obrigado